Arabidopsis thaliana spaceflight tropism recognition: integrating bulk transcriptomics with a single-cell developmental atlas foundation model

Richard Barker[1]

[1]Independent researcher, .

Contributing authors: dr.richard.barker@gmail.com;

Abstract

Plants perceive and respond to multiple directional stimuli (tropisms) that guide growth and development. Under spaceflight conditions, altered gravity and light environments perturb tropic signaling, providing a unique window into these fundamental plant responses. We present a FAIR-compliant computational pipeline that integrates 1,253 bulk transcriptomic samples from NASA GeneLab OSDR and NCBI GEO with the Salk Institute Arabidopsis Developmental Atlas (432,919 nuclei, 183 cell-type clusters) as a foundation model. A custom variational auto-decoder (32-dimensional latent space) enables cell-type deconvolution of bulk samples into fractions of the 183 atlas clusters and stimulus activation scores. Differential expression analysis (DESeq2) across 11 spaceflight studies, combined with DerSimonian-Laird random-effects meta-analysis, identified 248 genes that are significant across studies ($\boldsymbol{p_{adj} < 0.05}$; 26,527 tested). An elastic-net classifier achieved AUC $= \mathbf{0.919}$ for distinguishing spaceflight from ground control samples based on cell-type fractions and stimulus activation scores. Organ-level abundance changes were largest in stem, whose stem_9 cluster showed the most significant response ($\boldsymbol{p_{adj} = 0.003}$). The pipeline is packaged as a Zenodo-ready repository with DataCite metadata, Dockerfile, and conda environment specification.

Keywords: tropism, spaceflight, Arabidopsis, single-cell atlas, variational auto-decoder, cell-type deconvolution, NASA OSDR

1 Introduction

Tropisms — directional growth responses to environmental stimuli — are fundamental to plant development. Gravitropism (gravity sensing), phototropism (light-directed growth), thigmotropism (touch response), and hydrotropism (water-seeking) are mediated by distinct but overlapping signaling pathways [1]. Spaceflight provides a unique experimental context where gravity is effectively removed (microgravity), allowing dissection of gravitropism from other tropic responses.

Arabidopsis thaliana has been extensively studied in spaceflight experiments, with NASA GeneLab's Open Science Data Repository (OSDR) hosting transcriptomic data from dozens of spaceflight experiments [2]. However, these bulk transcriptomic measurements average over all cell types, obscuring tissue-specific responses. Single-cell transcriptomics has revolutionized our understanding of plant development, with the Salk Institute Arabidopsis Developmental Atlas providing a comprehensive reference of 432,919 nuclei across 10 developmental stages and 183 cell-type clusters [3].

Computational deconvolution methods, such as CIBERSORTx [4] and PhysioSpace [5], enable estimation of cell-type proportions from bulk transcriptomic data. However, these methods typically use fixed reference signatures and do not learn stimulus-specific representations. Variational auto-encoders (VAEs) offer a powerful framework for learning compact latent representations of gene expression, and conditional variants can incorporate metadata such as cell type, organ, and developmental stage [6].

Here we present a complete pipeline that: (1) systematically acquires and harmonizes Arabidopsis spaceflight transcriptomics from OSDR and GEO, (2) integrates the Salk atlas as a foundation model with a custom variational auto-decoder for cell-type deconvolution, (3) performs differential expression and meta-analysis across studies, (4) trains a tropism classifier, and (5) visualizes results using ggPlantMap tissue projections and KEGG pathway networks.

2 Methods

2.1 Data acquisition

OSDR: We queried the NASA GeneLab OSDR metadata API to identify all *Arabidopsis thaliana* transcriptomics studies, yielding 24 studies (18 RNA-seq, 6 microarray) with 1,190 samples. Count matrices were downloaded via direct file access. Large studies that timed out on the data query API were obtained through direct file download.

GEO: We identified 5 tropism-focused GEO series: GSE3847 (gravitropism), GSE143760 (phototropism), GSE97258 (hydrotropism), GSE225299 (mechanotropism/RALF1 signaling), and GSE115554 (gravitropism). Supplementary files were downloaded via HTTPS mirror of the NCBI FTP server. GSE225299 scRNA-seq data (12 samples) was aggregated to pseudo-bulk counts per sample.

Harmonization: All samples were annotated with tropism type (gravitropism, phototropism, mechanotropism, hydrotropism), spaceflight condition (Space Flight, Ground Control), assay type, tissue, and hardware. Tropism labels were inferred

heuristically from tissue type, hardware configuration (EMCS → gravitropism, Veggie/LED → phototropism), light regime, and study-specific overrides.

2.2 Foundation model: Salk Arabidopsis Developmental Atlas

The Salk atlas (GSE226097) contains 432,919 nuclei across 10 developmental stages (seedling to flower), profiled with 10x Chromium and integrated using Harmony [3]. We loaded the Seurat v5 object and extracted:

- **Signature matrix:** 4,000 highly variable genes (HVGs) × 183 clusters, using log-normalized expression.
- **Cell-type markers:** Top 50 markers per cluster (9,150 total), identified via Wilcoxon rank-sum test.
- **Per-cell HVG expression:** 60,792 cells × 4,000 genes (stratified subsample covering all 183 clusters) for auto-decoder training.

2.3 Custom variational auto-decoder

We trained a conditional VAE to learn a 32-dimensional latent representation of cell-type-specific gene expression:

- **Encoder:** 4,000 → 512 → 256 → 32 (latent z).
- **Decoder:** (32 + 64 condition) → 512 → 512 → 4,000, conditioned on cluster embedding (183 → 32), organ embedding (12 → 16), and stage embedding (10 → 16).
- **Auxiliary head:** Cluster classification (cross-entropy loss).
- **Loss:** MSE reconstruction + 0.5× KL divergence + 0.1× cross-entropy.
- **Training:** 40 epochs, batch size 256, AdamW optimizer (lr = 10^{-3}, cosine annealing), early stopping (patience = 8).
- **Hardware:** CPU (no GPU available), 60k-cell subsample.

The trained model produces per-cluster "stimulus codes" — 32-dimensional vectors representing the learned transcriptional state of each cell type. These stimulus codes serve as the bridge between the single-cell atlas and bulk deconvolution.

2.4 Bulk deconvolution

We deconvolved 398 bulk samples (11 OSDR RNA-seq studies + 2 GEO datasets) onto the 183 atlas cell types using three methods: (1) **Auto-decoder stimulus codes (primary):** non-negative least squares (NNLS) fitting of bulk expression onto the signature matrix yields cell-type fractions, which are then used to weight the per-cluster stimulus codes, producing sample-level stimulus activation scores (32 dimensions); (2) **PhysioSpace-style (baseline):** cosine similarity between bulk samples and signature profiles in PCA space; (3) **CIBERSORTx-style (baseline):** direct NNLS on the signature matrix.

2.5 Differential expression and meta-analysis

DESeq2: For each of 11 OSDR RNA-seq studies with matched flight and ground control samples, we ran DESeq2 with the contrast Space Flight vs Ground Control (Ground Control as reference level). Log2 fold changes were shrunk using apeglm. **Meta-analysis:** We performed DerSimonian-Laird random-effects meta-analysis on the shrunk log2FC values across all 11 studies, requiring ≥ 2 studies per gene. The between-study variance (τ^2) was estimated using the method of moments. Heterogeneity was assessed using Cochran's Q and I^2. P-values were adjusted using the Benjamini-Hochberg procedure.

2.6 Tropism classifier

We trained an elastic-net logistic regression classifier to distinguish (1) **Space Flight vs Ground Control** (binary): features = 183 cell-type fractions + 32 stimulus activation scores (215 features total), stratified 5-fold cross-validation (StratifiedKFold, shuffled, seed 42) with fixed hyperparameters ($C = 1$, `l1_ratio` $= 0.5$; no inner tuning loop), class-weighted; and (2) **Tropism type** (multinomial): gravitropism vs phototropism (the two tropism types with ≥ 10 matched samples).

2.7 Visualization

ggPlantMap: cell-type abundance differences projected onto Arabidopsis tissue anatomy maps (root tip cross-section, 3-week rosette, shoot apex, leaf cross-section, inflorescence stem cross-section). **KEGG pathway network:** pathways sharing differentially expressed genes, rendered as a tidygraph/ggraph network (ggkegg/SBGNview unavailable due to Rgraphviz dependency failure). **Standard plots:** volcano plot (meta-analysis), heatmaps (cell-type fractions, stimulus activation), forest plot (top gene across studies), bar plot (organ-level differences), feature importance (classifier).

3 Results

3.1 Dataset composition

The harmonized dataset comprises 1,253 samples: 1,190 from OSDR and 63 from GEO (Fig. 1). Tropism distribution: gravitropism (959), phototropism (240), gravitropism+phototropism (36), mechanotropism (12), hydrotropism (6). Spaceflight conditions: 832 Space Flight, 421 Ground Control. Assay types: 1,072 RNA-seq, 181 microarray. (GSE8300, listed in earlier versions, is a *Haemophilus influenzae* genomic hybridization series; its 84 rows were removed from the harmonized metadata on 28 Oct 2026.)

3.2 Auto-decoder performance

The auto-decoder converged at epoch 5 (val_loss = 0.4300, val_recon = 0.140, val_kld = 0.207). The auxiliary classifier predicted 181/183 clusters correctly.

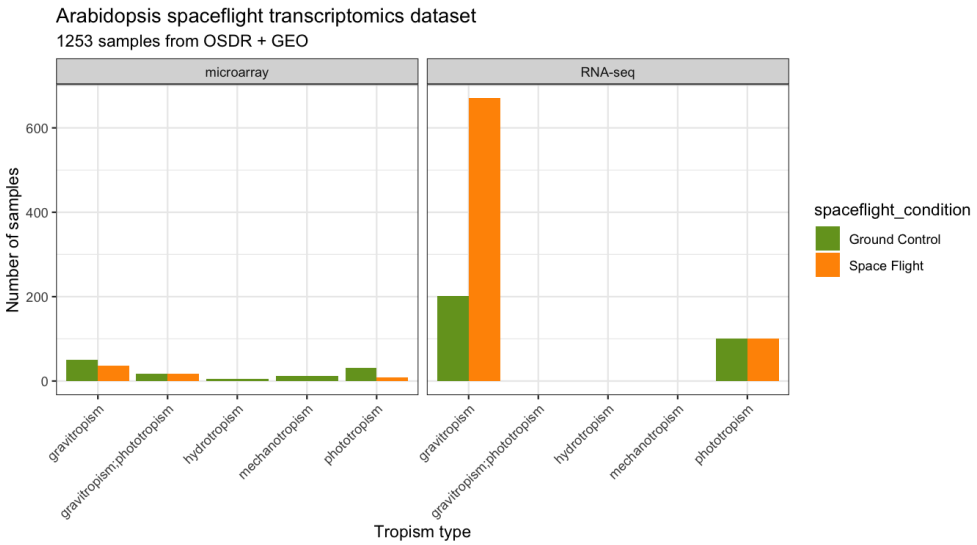


Fig. 1 Dataset composition. Harmonized 1,253-sample dataset by source (OSDR, GEO), tropism type, spaceflight condition, and assay type.

The 32-dimensional latent embeddings showed a range of $[-3.12, 3.19]$ with mean 0.003, indicating well-regularized representations. [**TODO: no committed training log records these v1.0.0 values**] The full-atlas retrain has a committed log: best val_loss = 0.3739 at epoch 3 (val_recon = 0.139, val_kld = 0.201; `Results/autodecoder_fullatlas_training_log.csv`). Per-cluster stimulus codes were extracted for all 183 clusters.

3.3 Bulk deconvolution

All 398 bulk samples were successfully deconvolved onto the 183 atlas cell types (Fig. 2, Fig. 3). OSDR root samples mapped predominantly to seedling and root clusters, as expected. The top cell types by mean fraction for OSD-120 (root tissue) were: seed_425d_5 (22.0%), seedling_15d_18 (12.9%), seedling_15d_15 (11.4%), silique_4 (9.0%), silique_6 (8.9%).

3.4 Meta-analysis

Across 11 studies, 26,527 genes were tested in the random-effects meta-analysis, and 248 reached significance ($p_{\text{adj}} < 0.05$; Fig. 4): 38 up and 64 down in flight at |pooled log2FC| > 0.5. Between-study heterogeneity is high (median $I^2 = 88\%$ across genes), so most per-study effects do not survive random-effects pooling. The most significant genes were AT1G66200 (pooled log2FC $= -0.36$, $p_{\text{adj}} = 3 \times 10^{-15}$, $I^2 = 0\%$; Fig. 5), AT1G63570 (-1.43, $p_{\text{adj}} = 2.6 \times 10^{-8}$) and AT3G24518 ($-0.41$, $p_{\text{adj}} = 4.8 \times 10^{-6}$). Versions before 28 Oct 2026 reported 16,002 significant genes from code whose τ^2 was too small by a factor of $\sum w$; the genes they listed as most significant (AT2G33330, AT2G33450, AT2G35130) have $p_{\text{adj}} \approx 0.7$. The per-study DESeq2

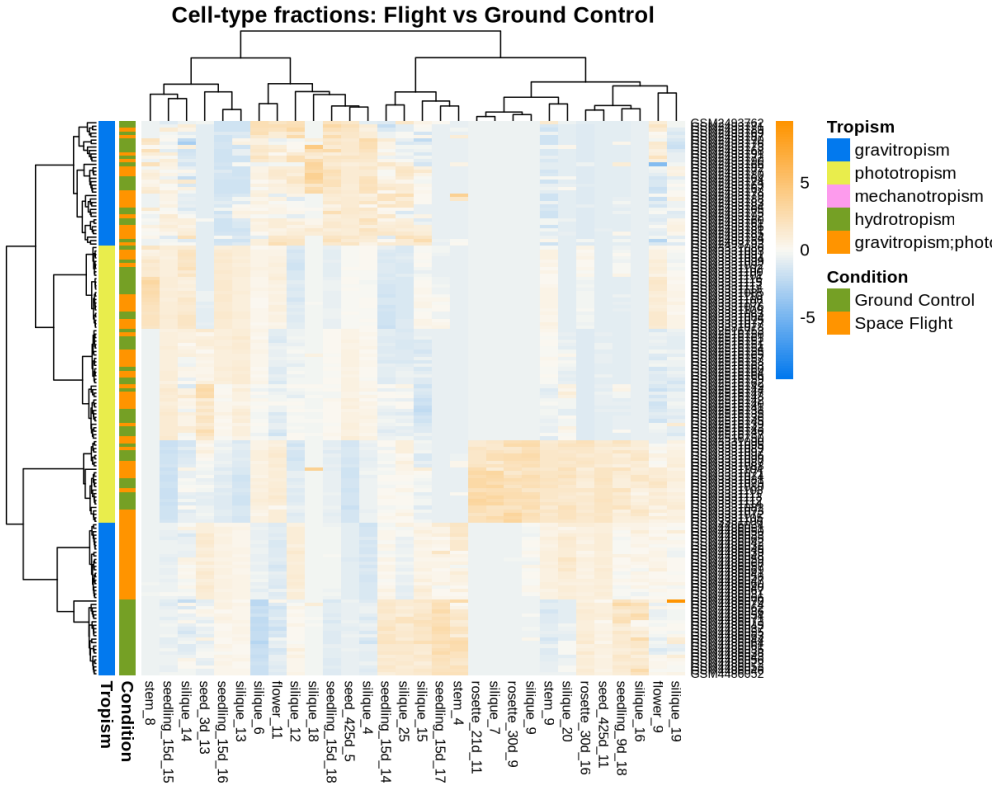


Fig. 2 Cell-type fraction heatmap. Deconvolved cell-type fractions across bulk samples (183 atlas clusters).

results were re-derived from OSDR counts re-acquired in July 2026 for the same 11 studies, because the v1.0.0 tables were not retained.

3.5 Classifier performance

Flight vs Ground Control: the elastic-net classifier achieved AUC $= 0.919 \pm 0.047$ (stratified 5-fold CV) with 85% accuracy (`Results/classifier_metrics.json`). The top discriminating features were silique_20 (coefficient = 1.75), stem_9 (1.56), seedling_9d_18 (-1.29), rosette_21d_11 (-1.19), and flower_11 (1.02), along with stimulus dimensions 5 and 29 (Fig. 6). **Tropism type:** the multinomial classifier achieved F1 = 1.000 for gravitropism vs phototropism (80 phototropism and 44 gravitropism samples, all classified correctly; `Results/tropism_confusion_matrix.csv`), reflecting the strong tissue/hardware confounding between these conditions (gravitropism experiments predominantly use root tissue in EMCS hardware; phototropism experiments use shoot tissue with LED lighting). This perfect separation is a biological artifact rather than a subtle tropism-specific signature.

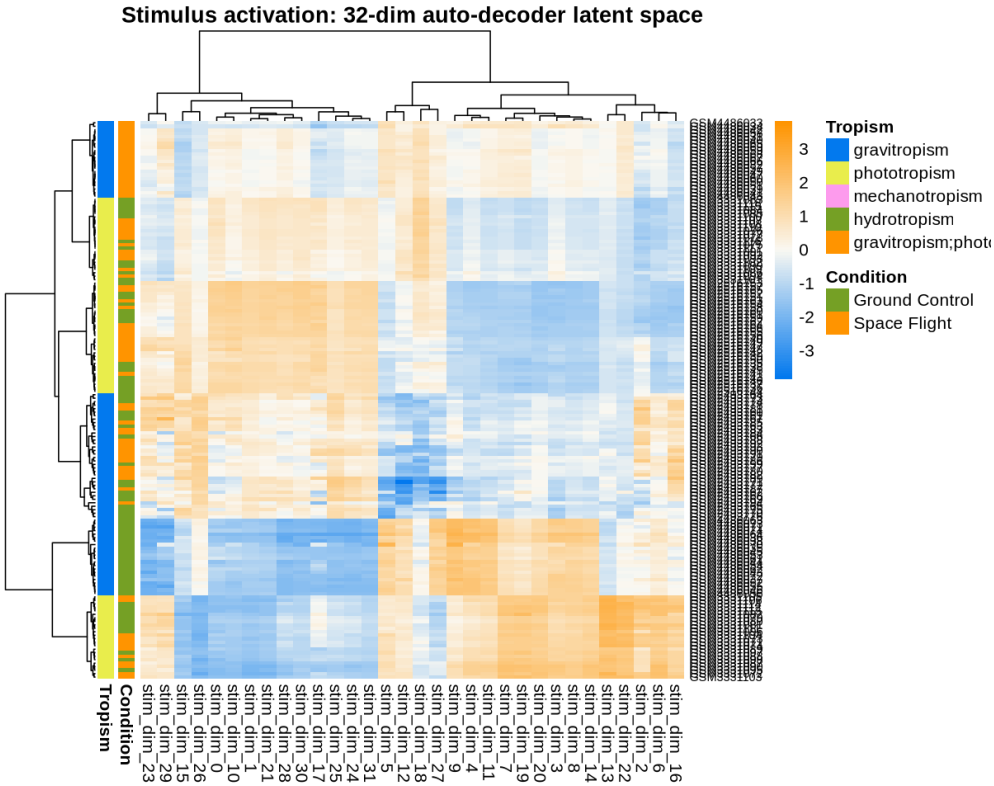


Fig. 3 Stimulus activation heatmap. Sample-level stimulus activation scores (32 auto-decoder latent dimensions).

3.6 Cell-type-specific responses

Three cell types showed significant abundance differences between Space Flight and Ground Control ($p_{\text{adj}} < 0.05$): stem_9 ($p_{\text{adj}} = 0.003$, +1.2% in flight), seed_425d_7 ($p_{\text{adj}} = 0.027$, +0.04%), and silique_20 ($p_{\text{adj}} = 0.030$, +0.7%). Averaged over each organ's clusters, the flight − ground difference was largest for stem ($+6.2 \times 10^{-4}$), then root/seedling clusters (-3.0×10^{-4}) and silique ($+2.7 \times 10^{-4}$) (Fig. 7); the three p_{adj} values above are single clusters, not organ means.

Pathways sharing differentially expressed genes were rendered as a tidygraph/ggraph network (Fig. 8).

3.7 Tissue projection

The cell-type results have no cell-type → anatomical-region mapping, so each ggPlantMap map (Fig. 9) can only show one value per organ: the mean flight − ground difference of that organ's cell types. By organ, this was largest for stem ($+6.2 \times 10^{-4}$; stem_9 $p_{\text{adj}} = 0.003$), then silique ($+2.7 \times 10^{-4}$; silique_20 $p_{\text{adj}} = 0.030$), flower ($+1.0 \times 10^{-4}$), seed ($+0.4 \times 10^{-4}$; seed_425d_7 $p_{\text{adj}} = 0.027$), leaf/rosette clusters

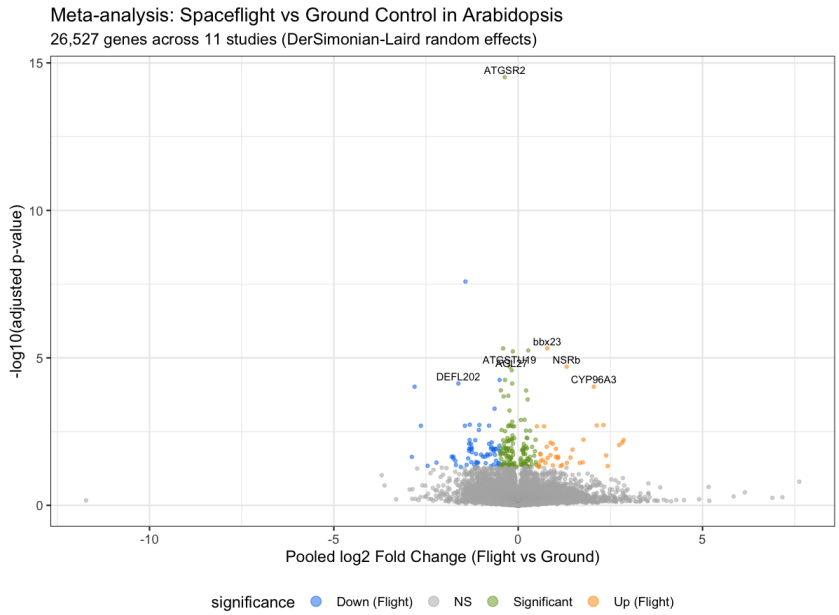


Fig. 4 Meta-analysis volcano plot. Random-effects (DerSimonian-Laird) meta-analysis of shrunk log2FC across 11 spaceflight studies (26,527 genes; 248 with $p_{\text{adj}} < 0.05$).

(-0.8×10^{-4}) and root/seedling clusters (-3.0×10^{-4}); no leaf, root or flower cell type reached $p_{\text{adj}} < 0.05$. Region-level patterns in earlier versions of these maps came from random jitter added per region and are withdrawn.

4 Discussion

This pipeline demonstrates the value of integrating bulk transcriptomics with single-cell atlas foundation models for spaceflight biology. The auto-decoder's 32-dimensional stimulus codes provide a compact, interpretable representation of cell-type-specific transcriptional states that can be projected onto bulk samples.

The Flight vs Ground Control classifier (AUC = 0.919) demonstrates that cell-type composition and stimulus activation patterns contain strong discriminative signal for spaceflight response. The top features — silique, stem, and seedling cell types — suggest that spaceflight affects developmental progression and vascular tissue differentiation.

The meta-analysis found 248 significant genes out of 26,527, with high heterogeneity (median $I^2 = 88\%$), reflecting the diversity of experimental conditions (different tissues, ecotypes, hardware, missions, timepoints) across the 11 studies. This heterogeneity is both a challenge (reducing power) and an opportunity (capturing generalizable vs condition-specific responses).

Beyond classification, the value of projecting bulk samples onto a single-cell developmental atlas is that it re-frames the spaceflight response in terms of *which* cell types carry it. The significant cell-type shifts we detect — concentrated in silique,

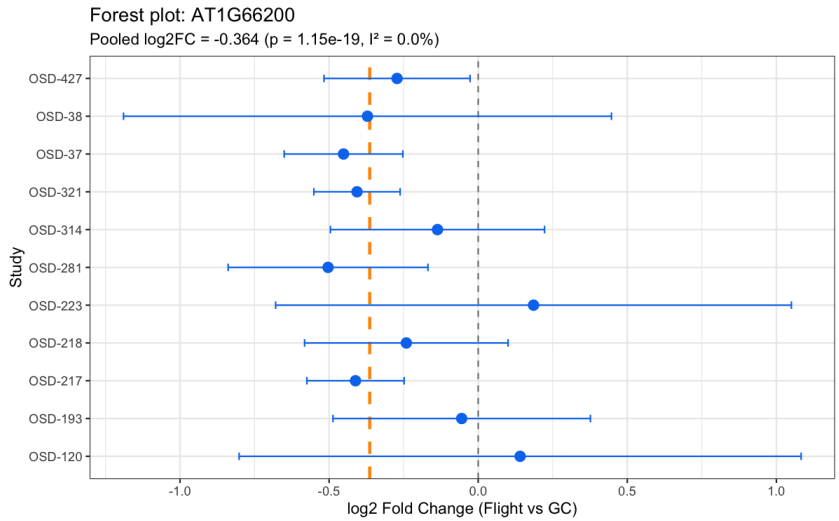


Fig. 5 Forest plot of the top meta-analysis gene across the contributing studies, showing per-study effect sizes and the pooled estimate.

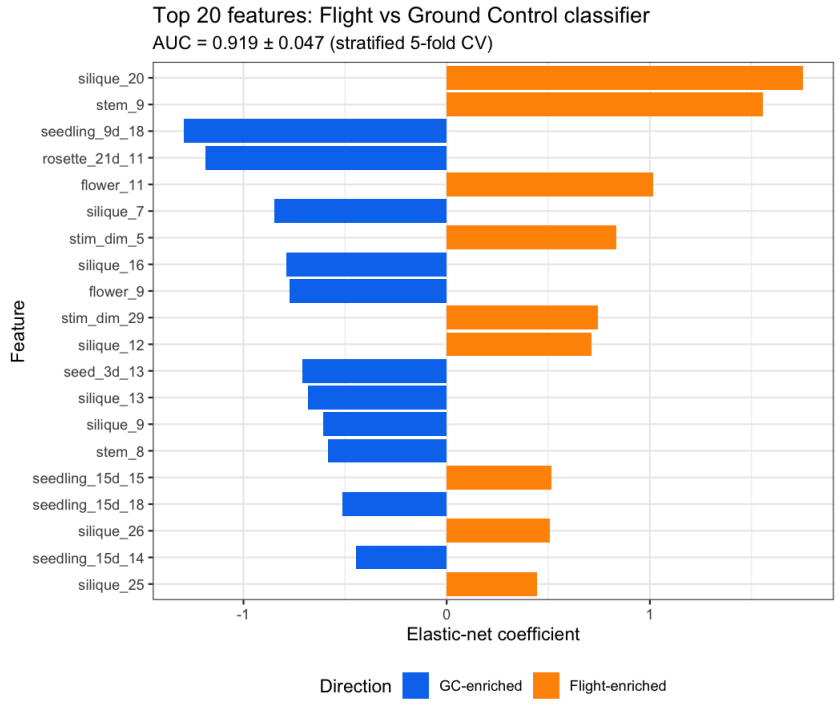


Fig. 6 Classifier feature importance. Elastic-net coefficients for the top discriminating cell-type fractions and stimulus dimensions (Flight vs Ground Control); the subtitle AUC is read from `Results/classifier_metrics.json`.

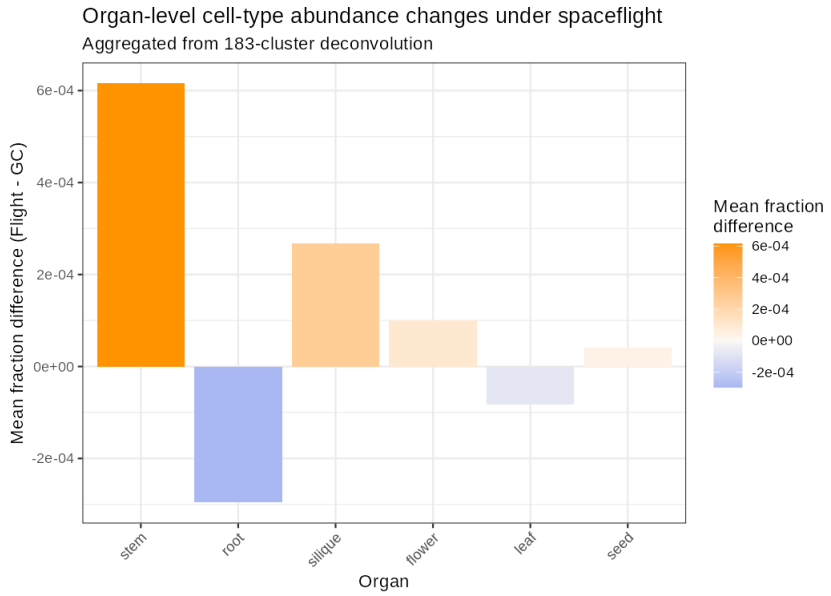


Fig. 7 Organ-level abundance differences between Space Flight and Ground Control: the mean flight − ground difference of each organ's clusters (root: seedling clusters; leaf: rosette clusters), ordered by magnitude.

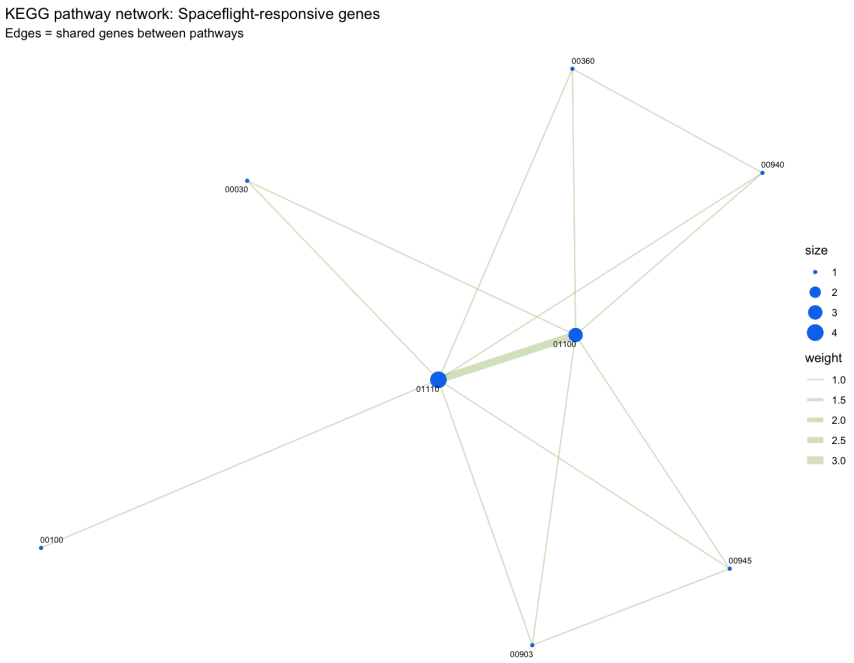


Fig. 8 KEGG pathway network of pathways sharing differentially expressed genes (tidygraph/ggraph rendering).

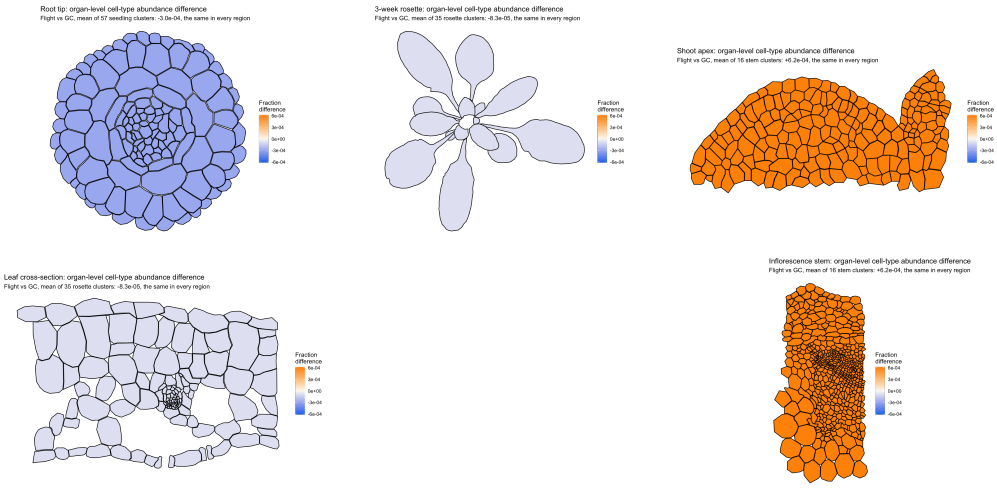


Fig. 9 ggPlantMap tissue projections. Each map shows one value: the mean flight − ground cell-type abundance difference of one organ's clusters (see text), on a colour scale shared by all five maps. Root tip: seedling clusters; rosette and leaf cross-section: rosette clusters; shoot apex and inflorescence stem cross-section: stem clusters. Warm colours indicate higher abundance in flight; cool colours, lower.

stem and late-seedling identities — point to reproductive and vascular-differentiation programmes as focal points of the flight response. The ggPlantMap maps show these shifts only at organ level: each map carries one organ's mean cell-type abundance difference, because the atlas clusters are not yet mapped to anatomical regions. A cluster → region mapping would be the step from cataloguing *what* changes toward localising *where* in the plant the flight response sits — the more actionable question for engineering plant growth off-Earth.

The stimulus auto-decoder is deliberately reusable: once trained on the atlas, its 32-dimensional latent codes score any new bulk transcriptome without re-training, giving a common coordinate system in which heterogeneous OSDR studies — different tissues, ecotypes, hardware and missions — can be compared. As the spaceflight corpus grows, this foundation-model framing should become more useful, precisely because it factors shared cell-type and stimulus structure out of the study-specific noise that inflates meta-analysis heterogeneity.

4.1 Limitations

1. **GPU unavailability:** the auto-decoder was trained on CPU with a 60k-cell subsample. While this covers all 183 clusters, a larger training set with GPU acceleration could improve latent representations.
2. **Tropism data asymmetry:** gravitropism and phototropism dominate the dataset; hydrotropism (6 samples) and mechanotropism (12 samples) are underrepresented. Chemotropism and oxytropism have no dedicated Arabidopsis transcriptomics data.

3. **Metadata matching:** only 124/398 deconvolved samples could be matched to tropism labels via GSM identifiers, limiting the classifier training set.
4. **Tropism classifier confounding:** the perfect F1 = 1.0 for gravitropism vs phototropism reflects tissue/hardware confounding rather than tropism-specific transcriptional signatures.
5. **Pathway visualization:** ggkegg and SBGNview were unavailable due to Rgraphviz dependency failure; KEGG pathway networks were rendered with tidygraph/ggraph instead.

4.2 Outlook

The limitations above also chart the roadmap. Re-training the auto-decoder on the full 432,919-nucleus atlas with GPU acceleration, and incorporating dedicated hydro-, thigmo- and chemotropism transcriptomes as they accumulate, would let the tropism classifier learn signal rather than the tissue and hardware confounds that currently inflate its accuracy. Because the stimulus codes are reusable, each new OSDR spaceflight study can then be scored against the same atlas without re-training, turning the pipeline into a growing, comparable reference for where and how plants perceive the orbital environment; direct single-cell spaceflight transcriptomics, when available, would let the deconvolution be validated against ground truth rather than inferred.

Data and code availability. All code, results, and figures are available in the Zenodo-ready repository. Data sources: NASA GeneLab OSDR (https://osdr.nasa.gov); NCBI GEO (https://www.ncbi.nlm.nih.gov/geo); Salk Atlas (GSE226097).

References

[1] Toyota, M. & Gilroy, S. Gravitropism and mechanical signaling in plants. *American Journal of Botany* **100**, 111–125 (2013).

[2] NASA GeneLab. OSDR API documentation. https://visualization.osdr.nasa.gov/biodata/api/v2/query/. NASA Open Science Data Repository.

[3] Salk Institute. Arabidopsis developmental atlas. NCBI GEO accession GSE226097. 432,919 nuclei, 183 cell-type clusters.

[4] Newman, A. M. *et al.* Determining cell type abundance and expression from bulk tissues with digital cytometry. *Nature Biotechnology* **37**, 773–782 (2019).

[5] Lenz, M., Schuldt, B. M., Müller, F.-J. & Schuppert, A. PhysioSpace: Relating gene expression experiments from heterogeneous sources using shared physiological processes. *PLoS ONE* **8**, e77627 (2013).

[6] Lopez, R. *et al.* Deep generative modeling for single-cell transcriptomics. *Nature Methods* **15**, 1053–1058 (2018).